

Agent Evaluation with MLFlow and EvalHub on Red Hat OpenShift AI

Version 1, 2026-08-11

Home

Goal

This course teaches how to deploy MLflow and EvalHub as included with Red Hat OpenShift AI 3.4; how to run evaluations of a self-hosted model, using EvalHub; how to run evaluations of an agentic application, using the MLflow SDK with Python; and how to visualize the results of those evaluations as MLflow experiments and traces using the Red Hat OpenShift AI dashboard.

Total estimated reading time: ...

You should reserve between 4 hours to one working day to complete this course.

This is the second in a series of quick courses covering MLflow, which consists of two courses:

1. [Agent Traces with MLflow and Red Hat OpenShift AI](#) (optional prerequisite course)
2. **Agent Evaluation with MLFlow and EvalHub on Red Hat OpenShift AI** (this course)



Red Hat associates must enroll using the [Red Hat Learning Portal](#), and Red Hat Partners must enroll using [Red Hat Partner Connect](#) to register this course in their learning history.

Objectives

On completing this course, you should be able to:

- Understand how MLflow and EvalHub relate to agentic applications and to Red Hat OpenShift AI
- Deploy a development environment with MLflow and EvalHub
- Evaluate LLMs using EvalHub
- Evaluate agents using MLflow

Audience

- Software developers working on agentic applications and agentic harnesses.
- Red Hat and Partner roles that perform demonstrations, PoC and PoV of MLflow with Red Hat OpenShift AI, such as Solutions Architects.
- Red Hat and Partner roles that implement and support MLflow and EvalHub with Red Hat OpenShift AI, such as Services Consultants, Technical Account Managers, and Customer Support Engineers.

Prerequisites

This course assumes that you have the following experience:

- Familiarity with generative AI and agentic application concepts.
- Familiarity with the basic operation of Red Hat OpenShift AI.
- Python programming skills are recommended but not required, because this training provides sample code which is ready to run.

It is recommended, but not required, to attend a prerequisite course: [Agent Traces with MLflow and Red Hat OpenShift AI](#).

Classroom environment

This course uses the [Red Hat OpenShift AI 3.4](#) demo catalog entry, which provides an SNO OpenShift cluster based on an AWS instance with a single GPU and a predeployed llama-32-3b-instruct model.

Any additional components required by MLflow, EvalHub, or course activities are configured by learners as part of course activities, using sample Python code and Kubernetes manifests provided in a public Git repository.

Other sources of information about MLflow and EvalHub

For documentation about Red Hat OpenShift AI, see its [product documentation](#), especially:

- [Working with MLflow](#)
- Chapter 2 of Evaluating AI systems: [Evaluate LLMs with EvalHub](#)

Also refer to upstream documentation of [MLflow](#) and [EvalHub](#).

Author

Fernando Lozano

Training Content Architect

Red Hat - Product Portfolio Marketing & Learning

Special thanks to Ahsen Shah, Edson Tirelli, and Matthew Prah for their support while designing and writing this course.

Lab Environment

Instructions to Launch Your Lab on the Red Hat Demo Platform (RHDP)

1. Log in to the [RHDP portal](#)
2. Search for the catalog entry: **Red Hat OpenShift AI 3**
3. Click the catalog entry in the search results.
4. On the catalog page, click **Order**.
5. Enter the required details in the order form and accept the defaults for other fields. You can use the smaller AWS instance type.
6. Review the warning at the bottom of the form and select:
I confirm that I understand the above warnings.
7. Await the email notification stating that your demo environment is ready and providing its access credentials and URLs.

Important Notes:

- This lab may take approximately **90** minutes or more to become ready.
- You can also retrieve access information directly from the RHDP portal.
- This lab uses GPU-enabled instances on AWS, which are in short supply. It may require multiple attempts to provision or restart.



Do **not** follow the instructions or showroom from the demo catalog entry. This course reuses that demo environment but does not require that you perform any of its original activities.

You will use the preprovisioned instance of Red Hat OpenShift AI as-is, and deploy MLflow, EvalHub, and sample agentic applications as part of this course activities.

How to Access Your Lab via the RHDP Portal:

1. On the RHDP portal, click **Services** option in the left-hand menu.
2. Select your lab from the services listings on the right-hand side of the page to view access details.

RHDP Portal Links

- Red Hat associates: <https://demo.redhat.com/>
- Red Hat partners: <https://partner.demo.redhat.com/>

Introduction to MLFlow and EvalHub

Goal

Understand how MLFlow and EvalHub relate to agentic applications and to Red Hat OpenShift AI.



Work In Progress

This chapter introduces MLflow and EvalHub in the context of Red Hat OpenShift AI, focusing on features relevant to generative AI and agentic applications. You will learn about the need to evaluate agentic applications and their components, such as large-language models (LLMs), and how it differs from traditional software engineering testing practices.

Evaluation Concepts for AI Applications

Objective

Describe how evaluating AI-enabled applications differs from testing traditional applications and identify the roles of MLFlow and EvalHub in Red Hat OpenShift AI.



Work In Progress

Evaluating non-deterministic applications

Traditional software engineering focuses on testing approaches for deterministic applications, such as unit tests and functional testings, and integrating that testing with other kinds of tests, such as performance tests and security scans.

AI-enabled applications, such as chat bots and agents, benefit from those traditional testing approach but they require the addition of tests designed for their non-deterministic nature. Those tests are called *evaluations*.

It is a known fact that large-language models (LLMs) are based on probabilistic pattern-matching. They may produce different outputs from the same input data. Sometimes those differences are inconsequential, such as a different wording of the same statements, but sometimes those differences make their answers factually incorrect or open to misleading interpretation.

Approaches such as guard rails can mitigate the probabilistic nature of LLMs, but cannot solve it entirely. Other approaches, such as setting the *temperature* parameter to zero, to force a model to always provide the same response, also degrades a model performance when planning, exploring, and researching. The randomness of LLMs actually enable them to work in imprecise scenarios and emulate creativity, problem-solving, or finding new insights. Without randomness, LLMs can only output facts which are already explicitly stated in their training data.

Model evaluations become popular quite early as a way of demonstrating the capabilities of LLMs and comparing them with each other based on criteria such as ability to provide correct answers to general knowledge questions, science questions, and other domain areas.

Other popular use cases for model evaluations include assessing how easy is to trick a model into providing a harmful response, or how capable they are at following instructions.

Evaluating models and applications

AI-enabled applications usually perform more than just sending a prompt to a model and displaying its output. Even chat bots are nowadays expected to function as *agentic applications* and rely on sophisticated *harnesses* which might include data such as:

- A *System prompt* with general instructions regarding the application's intended usage.
- Various *skills* which guide the model on performing specific tasks, and which are selected by the model based on the received prompt.
- Collections of *tools*, some provided by the agentic application itself, and some provided by external MCP servers.
- Data from *retrieval augmented generation (RAG)* sources, for example documentation of processes and legal or financial rules.
- Various levels of *memory* which add to the request information about the current user, past interactions, or the what's in the end user applications's screen.

Evaluating models, by themselves, is useful and necessary because the capabilities and performance of an LLM vary with their target hardware and configuration settings. You should assess that, on your specific environment, a model performs as expected regarding correctness of its outputs, resistance to bad prompts, frequency of hallucinations, and tool capability.

But evaluating a model is not sufficient to assess the performance of agentic applications. You must evaluate individual components of its harness, such as their RAG sources and MCP servers. You must also evaluate the harness as a whole and experiment with different combinations, for example:

- Among different MCP servers that integrate with the same back-end systems, the most powerful MCP server may perform worse when combined with other elements of a harness.
- A simpler MCP server, which performs worst individually, may be better when combined with a specific system prompt or set of skills.
- Preloading data from memory and RAG may be better or worse than letting the model request such data on-demand from MCP servers.

Most applications require multiple evaluations, of different kinds, to assess a good enough performance and quality of their responses. They also require repeated evaluations to catch regressions when any piece of their harnesses change, and their evaluations sets will grow over time, as you discover different usage scenarios and failure modes.

Evaluations during development and production

Evaluations are conceptually closer to performance tests than to unit tests. You cannot just test one example of "good" input plus one example of each known corner case. You must test multiple "good" inputs and also multiple examples of each corner case, because of the non-deterministic nature of LLMs. You may require large evaluation data sets and also curate custom data sets which

are specific for your usage scenarios and business processes.

Besides, your evaluations are constrained by the availability of specialized hardware such as GPUs and specialized memory for storing models and tokens. You will not run them for every commit.

Early during development, your evaluations will be exploratory and interactive, intended to select an optimal harness, and use a small data set. Once your harness and model selections are more solid, you will increase the size of your evaluation data sets and add a larger variety of evaluations to assess performance, correctness, and the need for specialized guardrails. You will also start to run evaluations automatically, as part of your CI/CD processes.

Once you release your agent to production, or to its first external test users, you will collect samples (traces) of actual inputs and outputs to assess actual performance of your harness, and feedback some of that data to improve the evaluations for the next development interaction of your agent.

It is recommended that you work proactively: instead of relying just on user feedback about the quality of your application, such as thumbs up and down, and customer support tickets complaining about errors, you should take samples of actual usage data as additional evaluation data sets and compare the results with evaluations from your synthetic and development evaluation data sets.

You may learn that actual behavior differs significantly from intended and tested behavior, and consequently improve both your harness and its evaluations.

Running evaluations with MLflow

MLflow provides fundamental observability and experiment-tracking capabilities that enable developers and data scientists to ensure their applications work beyond a demonstration and perform as expected in production.

MLflow is a popular library for agentic evaluations. It requires encapsulating your agent and its harness as a *predictor function*, and then creating custom code to invoke your prediction function together with *scorers* and *judges*, some of which are themselves based on LLMs.

It is up to the developer packaging their custom evaluation code, based on the MLflow SDK, as a KubeFlow pipeline or as a Tekton task (for OpenShift Pipelines), among other ways of operationalizing reproducible execution of MLflow evaluations.

MLflow is also a popular library for tracing agentic applications, including their usage of RAG sources, tools, and multi-turn conversations with LLMs.

Typical usage of MLflow libraries store evaluation results as part of a trace, and multiple traces are part of an *experiment*. An MLflow experiment could include other kinds of data and other kinds of experiment runs besides traces.

Using the MLflow SDK, you can also retrieve data from traces as evaluation data sets for running evaluations, closing the loop from development to production, or from early development tests to test users.

Running evaluations with EvalHub

Most kinds of AI evaluations are run by custom code, using varied libraries which specialize in different kinds of evaluations, such as safety and domain expertise. As you can guess, maintaining and operationalizing such custom code quickly becomes a challenge.

EvalHub is an open-source evaluation orchestration platform designed to run systematic, multi-framework tests for large language models (LLMs), agents, and AI systems.

EvalHub standardizes how you operationalize and store results of your evaluations. It provides a number of *evaluation providers*, each specific for an *evaluation framework* and their specific *benchmarks*. EvalHub also offers a standard syntax to combine multiple providers and select a subset of their benchmarks in a *collection*, which makes it easier to reproduce evaluation runs.

EvalHub also improves governance of your evaluations by storing their results on OCI container images, for immutable auditability storage.

As included with RHOAI 3.4, EvalHub supports a number of providers focused on evaluating models, such as Garak and LightEval. The EvalHub community is very active and it is expanding its set of community-supported providers, which add capabilities for running other kinds of evaluations, including agent evaluations. As these community providers mature, some of them will be included and supported with future releases of RHOAI.

Comparing and integrating EvalHub with MLflow

It is not a choice of using either EvalHub or MLflow for your evaluations. Most of the times, it is about how you will use each of them, and sometimes them together, during your AI application lifecycle.

While MLflow provides building blocks for creating custom evaluations for your agents and its business needs, EvalHub provides standard benchmark collections which are ready to run and compare how different models perform on your specific deployment environment. EvalHub also enables customizing those collections, by selecting different benchmarks or the weights and thresholds from each individual benchmark to compose the overall evaluation score.

RHOAI uses MLflow as its main observability feature. Multiple components of RHOAI integrate with MLflow to store traces and metrics as MLflow experiments. Among those components, EvalHub can store evaluation results from any provider and benchmark as MLflow experiments. This way, results from all your evaluations, be them managed by EvalHub or not, are stored and visualized together in MLflow.

Unfortunately, at the time of RHOAI 3.4, EvalHub is not yet capable of running MLflow evaluations directly. It would require building custom code using the bring your own framework (BYOF) API from EvalHub to run an MLflow evaluation.

This course does not present BYOF, but provides you with sufficient information so you could use EvalHub for run model evaluations and MLflow for agent evaluations, and combine both as discrete steps in a KubeFlow pipeline, while aggregating results of all these evaluations in MLflow.

What's Next

Now that you understand the key evaluation concepts for AI applications, the next activity explores the hands-on lab environment and the next chapter deploys MLflow and EvalHub on Red Hat OpenShift AI.

Lab: Explore the Learning Environment

Objective

Navigate the provided demo environment, identify enabled features of Red Hat OpenShift AI 3.4, and verify the predeployed llama model.



Pending Review

Before you begin

You need access to the predeployed Red Hat OpenShift AI (RHOAI) instance provided in your lab environment. Ensure you have:

- Access credentials for the RHOAI dashboard
- A web browser to access the RHOAI dashboard

Instructions

In this lab, you will log in to the RHOAI dashboard and explore its main features and components to understand the platform's capabilities.

Inspect the existing projects and deployed models

1. Access the RHOAI Dashboard.
 - a. Navigate to the RHOAI dashboard URL of your demo environment.
 - b. Log in using the **admin** user, from the credentials provided for your demo environment.
2. Explore the precreated project.

Your RHOAI instance contains a model which is already deployed in a regular project named **my-first-model** — which is **not** an AI project.

That model takes up all GPU capacity available to your RHOAI instance, so any change to the model deployment requires that you stop it and restart after making changes.

- a. In the left pane, select **AI hub** > **Models** and then select the **Deployments** tab.
- b. In the **Project** field, enter **my** and select the **my-first-model** project.



Depending on the current page of the RHOAI dashboard, it may be necessary to first select **All projects** before you can search and select the **my-first-model** project.

- c. Verify that there is a model deployment named **llama-32-3b-instruct** and its status is **Ready**.
 - d. Click **View**, under the **Endpoints** column for the model row, and verify that it provides only an internal endpoint. This means the model cannot be accessed from outside of its OpenShift cluster, unless you change its model deployment.
3. Try the model with a Gen AI studio playground.
 - a. In the left pane, select **Gen AI studio › Playground**.
 - b. In the **Project** field, enter **my** and select the **my-first-model** project.



The project previously selected in one page of the RHOAI dashboard may not be the current selection in another page of the RHOAI dashboard.

- c. If you do not see a **Configure** pane, click **Settings**. In the **Configure** pane, select the **Model** tab and verify that the selected model is named **llama-32-3b-instruct**. Ignore the prefix on the model name.
- d. In the right pane, click **Send a message...** and type a prompt, for example:

What is RHEL?

- e. Type **Enter** to submit the prompt and see the response from the model. The model may return different answers if you attempt the same prompt again, but they all should be correct and reasonable.

Confirm that MLflow and EvalHub are not deployed in RHOAI

4. Verify that there are no RHOAI dashboard pages related to MLflow.
 - a. In the toolbar, to the top right, click the **Applications** icon, which is next to the **Question** icon, and verify that there is **not** an entry named **MLflow**.
 - b. In the left pane, expand the **Develop & train** category, which contains entries named **Feature store**, **Pipelines**, **Jobs**, and **Evaluations**.
 - c. In the left pane, expand the **Gen AI studio** category, which contains entries named **Playground** and **AI asset endpoints**.

Enabling MLflow will add additional entries here.

5. Verify that there are no MCP servers configured in the GenAI Studio.

Select **Gen AI studio › AI asset endpoints** and then select the **MCP servers** tab.

The page should list no MCP servers.

6. Verify that the Evaluations page from RHOAI dashboard is not functional.
 - a. In the left pane select the **Develop & train › Evaluations**.
 - b. The **Evaluations** page display the messages **Evaluations unavailable** and **To use evaluations, enable the evaluation service using the Trusty AI operator**.

Summary

After completing these steps, you verified that you have a functional RHOAI instance with a pre-deployed model, which is available only for internal access, and that you can submit prompts to the model using an RHOAI playground.

In addition, you have also verified that MLflow and EvalHub are is not installed on your RHOAI instance, and that there are no MCP servers deployed.

What's next

Now that you are familiar with the learning environment, the next chapter covers how to deploy MLflow and EvalHub for development and testing usage on Red Hat OpenShift AI.

Deploy MLFlow and EvalHub

Goal

Deploy a development environment with MLFlow and EvalHub on Red Hat OpenShift AI.



Work In Progress

This chapter describes how to deploy MLflow and EvalHub using operators from RHOAI. It also describes the architecture of MLflow and EvalHub and how they integrate with the authentication and authorization model of RHOAI and OpenShift.

MLFlow Architecture and Integration

Objective

Describe the architecture of MLFlow and explain how to enable and access it in Red Hat OpenShift AI.

Architecture of MLflow

MLflow is based on a client/server architecture. An AI application, which is instrumented by using either the MLflow SDK or OpenTelemetry, sends observability data, such as metrics and traces, to an MLflow server, which stores that data and provides HTML-based visualization of it.




Figure 1. Architecture of MLflow

It does not matter whether the "application" is a user-facing application, an autonomous agent, or a training or evaluation pipeline. It just needs to send observability data to an MLflow server.

The MLflow server makes use of a relational database and an artifact store to record artifacts such as prompts and responses. It is recommended that the artifact store be an S3 service, for scalability.

The relational database stores metadata and tags which enables the MLflow server to relate all inputs, outputs, and observability data from a single conversation with an LLM to each other. PostgreSQL is a popular option for the MLflow database.

MLflow experiments and workspaces

Experiments are a key concept in MLflow. An *experiment* is a grouping of observability data from multiple runs. Each run collects all data obtained while running an AI application once. Runs can take multiple forms, such as *traces*, which record the interactions between an agent and a model.

A *workspace* is a grouping of experiments. Typically, a workspace corresponds to an application, and experiments correspond to running the application with different configuration parameters, such as different models, or different system prompts.

Then, inside an experiment, each run connects an input to its output, and provides metrics and other observability data from that run.

The MLflow Operator and RHOAI

There is an MLflow operator included with Red Hat OpenShift AI (RHOAI), but this operator will not appear in the software catalog or in the list of installed operators managed by the Operator Lifecycle Manager (OLM). The MLflow operator is managed by the RHOAI operator itself.

To deploy an MLflow server, you must perform two actions:

1. Edit the singleton data science cluster resource of your RHOAI instance to enable the MLflow operator by setting its `spec.components.mlflowoperator.managementState` field to `Managed`.
2. Create an MLflow resource, which configures the database and artifact storage of the MLflow server.

Once the MLflow resource is available, the RHOAI dashboard displays a number of new pages and adds more elements to other pages. The **Develop & train › Experiments** page is the main user interface to MLflow in RHOAI 3.4.

MLflow workspaces and authentication within OpenShift

The MLflow operator maps Kubernetes namespaces and OpenShift projects (not just AI projects) to MLflow workspaces. For each project in RHOAI, there is an MLflow workspace.

The MLflow operator also configures the MLflow server to delegate authentication and authorization to OpenShift. Any OpenShift user is a valid MLflow user. Such users have access to data in MLflow workspaces based on their RBAC permissions from the respective OpenShift project.

Service accounts also get access to the MLflow workspace, so pods running in OpenShift can access the MLflow server in RHOAI. These pods include GenAI studio playgrounds and Jupyter Labs notebooks.

If you must connect to MLflow from an application running outside of its OpenShift cluster, you can take the same user token you would use for access to Kubernetes APIs and use it to authenticate to

the MLflow server. External applications connect to the MLflow server using the same route (URL) a web browser uses to access the standalone MLflow web UI.

What's Next

Next, you deploy and configure MLflow on your Red Hat OpenShift AI instance.

Lab: Deploy MLflow for Development

Objective

Deploy and verify a development instance of MLflow in Red Hat OpenShift AI.



Work In Progress

Before you begin

You need access to the predeployed Red Hat OpenShift AI (RHOAI) instance provided in your lab environment. Ensure you have:

- Access credentials for the RHOAI dashboard
- A web browser to access the RHOAI dashboard
- Access credentials for the OpenShift web console

If you are in a hurry and need to quickly enable MLflow in your demo instance, for example: you had to destroy and recreate it before completing all activities in this course, you can:



- Clone the [samples git repository](#)
- Enter its `scripts` directory
- Log in to OpenShift in your demo instance as the `admin` user, using the `oc` command
- Invoke the `enable-mlflow.sh` script

Instructions

You will enable MLflow on your Data Science Cluster singleton resource and deploy a development instance of the MLflow server. Then you will create a Jupyter notebook to test connectivity to the MLflow server. You will use the MLflow SDK for Python.

As you perform this activity, you will alternate between three browser tabs:

- The Red Hat OpenShift web console
- The Red Hat OpenShift AI (RHOAI) dashboard
- A Jupyter notebook from a RHOAI workbench

Once you open one of these tabs, keep it open because you may return to it later.

Install and configure the MLflow Server

1. Enable MLflow on RHOAI by editing its data science cluster resource.
 - a. Open the OpenShift web console and log in as a user with **cluster-admin** rights.
 - b. In the left navigation pane, select **Ecosystem > Installed Operators**. Then click **Red Hat OpenShift AI**.
 - c. In the Red Hat OpenShift AI operator page, select the **Data Science Cluster** tab. Then click the **default-dsc** resource.
 - d. In the **default-dsc** resource page, select the **YAML** tab. In the text editor, search for the **spec.components** field, and set **mlflowoperator.managementState** to **Managed**.

```
...
spec:
  components:
    mlflowoperator:
      managementState: Managed
    kserve:
      managementState: Managed
      rawDeploymentServiceConfig: Headless
...
```

- e. Click **Save** and wait until you see a confirmation message similar to:

```
alert:default-dsc has been updated to version 1186240
```

- f. If you receive an error message, fix the errors in your YAML and try saving again.
2. Configure your MLflow instance by creating a MLflow resource.
 - a. Open the [sample MLflow manifest](#) from the course samples repository. Copy the entire contents of the **scripts/mlflow.yaml** manifest to the clipboard.



That manifest corresponds to the example from [Minimal development or test deployment](#) in the Red Hat OpenShift AI product documentation with no changes.

The full manifest is listed here, just for your reference:

```
apiVersion: mlflow.opendatahub.io/v1
kind: MLflow
metadata:
  name: mlflow
  namespace: redhat-ods-applications
spec:
```

```
storage:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
backendStoreUri: "sqlite:///mlflow/mlflow.db"
artifactsDestination: "file:///mlflow/artifacts"
serveArtifacts: true
```

- b. In the left navigation pane, select **Home** > **API Explorer**. Then enter **MLflow** in the search field. You will see three resource types; click **MLflow**.
- c. In the MLflow Resource details page, select the **Instances** tab and click **Create MLflow**.
- d. In the YAML text editor, delete all of its sample contents and paste the manifest you already copied to the clipboard.
- e. Click **Create** and, in the MLflow details page, wait until the **Available** condition is **True** and the **Progressing** condition is **False**.

Test your MLflow server using the dashboard and a workbench

3. Verify MLflow is available from RHOAI and create an empty experiment.
 - a. Open your RHOAI dashboard. If it is already open, refresh your browser window.
 - b. In the left navigation pane, expand **Develop & train**. Notice the new item named **Experiments (mlflow)** and click it.
 - c. In the **Project** field, select the **my-first-model** project.
 - d. Click **Create Experiment** and enter **dummy experiment** as its name.
 - e. Verify that you see the Experiments page for **dummy experiment** but there is no data on any of its pages, such as **Traces**, **Sessions**, or **Prompts**.
 - f. Creating an empty experiment is sufficient to prove that your MLflow instance can connect to its database and artifacts storage.
 - g. You can, alternatively, open the upstream MLflow web interface by clicking the **applications** menu, near the top right of the navigation bar close to the help icon, and selecting **MLflow**. You should not need it: all functionality you need for this course is available from the RHOAI dashboard.
4. Test connectivity to MLflow from a Jupyter notebook inside a workbench.



If you have trouble typing the Python code in this and other activities of this course, or you cannot cut-and-paste from the course page, you can access the complete source code for all Jupyter notebooks in the [course samples git repository](#).

- a. In the left navigation pane, click **Projects** and select **All projects** in the **Project** field. Then enter **my** in the **Filter by name** field and click **my-first-model** in the list of projects.

- b. In the Projects page for `my-first-model`, select the **Workbenches** tab and click **Create workbench**.
- c. In the Create workbench page, enter `test-evals` for the **Name** field and, in the **Workbench image** field, select **Jupyter | Minimal | CPU**.
- d. Leave all other fields with their default values and click **Create workbench**.
- e. In the Workbenches list, wait until the row for the `test-mlflow` workbench displays a **Ready** state and then click **test-mlflow** to open its Jupyter Lab in a new browser tab.
- f. Click the **Python 3.12** Notebook icon to create an empty notebook. In the left navigation pane of the Jupyter Lab, right-click the **Untitled.ipynb** file, which is your new notebook, and select **Rename**. Then, enter `test-mlflow.ipynb` as its new file name.
- g. In the first cell of your notebook, enter:

```
! pip install mlflow[kubernetes]
```

- h. In the top toolbar from the notebook, click the **play** icon, having the tooltip "Run this cell and advance".
- i. Wait until the MLflow SDK and all its dependencies are loaded from the RHOAI index into the workbench.
- j. After the `pip` command finishes, enter the following on the second cell of your notebook:

```
import os
import mlflow

os.environ["MLFLOW_TRACKING_AUTH"]="kubernetes-namespaced"
mlflow.set_tracking_uri("https://mlflow.redhat-ods-
applications.svc.cluster.local:8443")

print(mlflow.list_workspaces())
```



In this example, we are using the *internal* URL of the MLflow server, which works only for applications running in the same OpenShift cluster.

- k. In the top toolbar from the notebook, click the **play** icon. The results should be an array containing a single **Workspace** object named `my-first-model`.

This notebook takes advantage of the integration between the RHOAI dashboard and its managed instance of the MLflow server. It connects using the service account of the `test-evals` workbench, which is authorized to access the MLflow workspace corresponding to its `my-first-model` project.

- l. You will not use the `test-mlflow.ipynb` notebook anymore, but do NOT delete the `test-evals` workbench because you will keep using it in the remaining activities of this course.

Summary

After completing these steps, you have verified that you have a minimal but functional MLflow instance, to which you can connect from applications running inside your RHOAI cluster and store data from your experiments.

What's next

With MLflow deployed and verified, the next section introduces the architecture and setup of EvalHub.

EvalHub Architecture and CLI

Objective

Describe the architecture of EvalHub and explain how to enable, access, and use its CLI with Red Hat OpenShift AI.



Work In Progress

Architecture of EvalHub

The architecture of EvalHub is similar to that of Kubernetes controllers and operators in the sense that it manages and observes evaluation runs which are ultimately implemented as Kubernetes jobs.




Figure 2. Architecture of EvalHub

EvalHub offers its own API for managing and querying the state of evaluation runs instead of using Kubernetes custom resources. This makes EvalHub easier for AI engineers and software developers not experienced with Kubernetes. It also enables the community to deploy and use EvalHub by itself, outside of Kubernetes.



Standalone deployments of EvalHub are *not* supported by Red Hat OpenShift AI.

A client uses the EvalHub API to submit requests to create and query the state of evaluation runs, evaluation providers, and collections. The RHOAI dashboard acts as an EvalHub client, and there is an EvalHub CLI tool included with the EvalHub SDK for Python. Among other potential clients for the EvalHub server are Jupyter notebooks and KubeFlow pipelines.

The EvalHub server uses a PostgreSQL database to store the state of evaluation runs, including their metrics and pointers to the optional MLflow experiment and OCI container image which also store their evaluation results.

EvalHub and TrustyAI in RHOAI

There is no EvalHub operator in RHOAI. The TrustyAI operator manages EvalHub servers in RHOAI. The TrustyAI operator, in turn, is managed by the RHOAI operator. Consequently, the TrustyAI is not visible from the Operator Lifecycle Manager (OLM).

To deploy an EvalHub server, you must perform two actions:

1. Edit the singleton data science cluster resource of your RHOAI instance to enable the TrustyAI operator by setting its `spec.components.trustyai.managementState` field to `Managed`.
2. Create an EvalHub resource, which configures the database, MLflow server, and optional OCI container registry associated with the EvalHub server.

Once the EvalHub resource is available, the **Develop & train > Evaluations** page of the RHOAI dashboard becomes functional and exhibits the **Create new evaluation** button.



Contrary to the note on the RHOAI 3.4 product docs, you *must* create your EvalHub resource in the `redhat-ods-applications` namespace to enable its RHOAI dashboard integration.

If you create your EvalHub resource in any other namespace, you cannot use the RHOAI dashboard to manage evaluation runs, but you can still use the EvalHub SDK and its CLI.

If you install EvalHub in the standard namespace for RHOAI applications, as recommended, you must set the TrustyAI label in any project where you run evaluations:

```
evalhub.trustyai.opendatahub.io/tenant
```

Without that label, network policies would prevent pods in the project, including pods from Kubernetes jobs created by EvalHub itself to run evaluations, from connecting to the EvalHub server. It doesn't matter the value assigned to the label. Only its key is necessary.

EvalHub tenants and authentication within OpenShift

The EvalHub operator maps Kubernetes namespaces and OpenShift projects (not just AI projects) to EvalHub tenants. For each project in RHOAI, there is an EvalHub tenant.

The EvalHub operator also configures the EvalHub server to delegate authentication and authorization to OpenShift. Any OpenShift user is a valid EvalHub user. Such users have access to EvalHub tenants and their evaluation runs based on their RBAC permissions from the respective OpenShift project.

Service accounts also get access to EvalHub, so pods running in OpenShift can access the EvalHub server in RHOAI. These pods include Jupyter Labs and KubeFlow pipeline tasks.

If you must connect to EvalHub from an application running outside of its OpenShift cluster, for example the EvalHub CLI, you can take the same user token you would use for access to Kubernetes APIs and use it to authenticate to the EvalHub server.

Installing the EvalHub CLI

The EvalHub CLI is provided as part of the EvalHub SDK, which is a Python library. The recommended approach is creating a Python virtual environment and installing the EvalHub SDK and CLI inside it.

Once you have the `evalhub` command available on your shell, you must use the `evalhub configure`

command to set the URL of the EvalHub server, the tenant, and authentication token so you can create and manage evaluation runs.

What's Next

Next, you deploy and configure EvalHub on your Red Hat OpenShift AI instance.

Lab: Deploy EvalHub for Development

Objective

Enable EvalHub on Red Hat OpenShift AI using the TrustyAI operator, configure its database, and install and configure the EvalHub CLI.



Pending review

Before you begin

You need access to the predeployed Red Hat OpenShift AI (RHOAI) instance provided in your lab environment. Ensure you have:

- Access credentials for the RHOAI dashboard
- A web browser to access the RHOAI dashboard
- Access credentials for the OpenShift web console

Also ensure that you already deployed MLflow by following the instructions from the [previous lab](#) because, if you enable EvalHub without having MLflow enabled, then EvalHub will not be able to save metrics on experiments from MLflow.



If you already deployed MLflow, and you are in a hurry and need to quickly enable EvalHub in your demo instance, for example: you had to destroy and recreate it before completing all activities in this course, you can:

- Clone the [samples git repository](#)
- Enter its `scripts` directory
- Log in to OpenShift in your demo instance as the `admin` user, using the `oc` command
- Invoke the `enable-evalhub.sh` script

Instructions

You will enable TrustyAI on your Data Science Cluster singleton resource and deploy a development instance of the EvalHub server. Then you will use a shell to test connectivity and healthy of your EvalHub server.

As you perform this activity, you will alternate between three browser tabs:

- The Red Hat OpenShift web console
- The Red Hat OpenShift AI (RHOAI) dashboard

Once you open one of these tabs, keep it open because you may return to it later.

Install and configure the EvalHub Server

1. Verify that TrustyAI is already enabled in the data science cluster resource.
 - a. Open the OpenShift web console and log in as a user with `cluster-admin` rights.
 - b. In the left navigation pane, select **Ecosystem** > **Installed Operators**. Then click **Red Hat OpenShift AI**.
 - c. In the Red Hat OpenShift AI operator page, select the **Data Science Cluster** tab. Then click the **default-dsc** resource.
 - d. In the `default-dsc` resource page, select the **YAML** tab. In the text editor, search for the `spec.components` field, and verify that `trustyai.managementState` is already set to **Managed**.

```
...
spec:
  components:
    kserve:
      managementState: Managed
      rawDeploymentServiceConfig: Headless
  ...
  trustyai:
    managementState: Managed
    mcpGuardrailsMode: false
  aipipelines:
    managementState: Managed
  ...
```

- e. Click **Cancel** to leave the YAML editor without making changes.
2. Deploy a PostgreSQL database.
 - a. Open the [sample evalhub db deployment](#) from the course samples repository with a web browser and open the manifest `scripts/evalhubdb.yaml`. Copy the entire contents of the manifest to the clipboard.



This manifest provides a minimal PostgreSQL deployment. It is not listed here because of its size, but you would not use such a simplistic database deployment for a real development or production instance of EvalHub.

You are free to use your preferred OpenShift template, helm chart, or database operator instead. Ideally, use something that is supported by your organization IT department, and for which you can expect product support from a vendor.

- b.

In the left pane, select **Home** > **Projects** and, in the search field, enter `redhat-ods`. Then click the **redhat-ods-applications** project.

- c. In the toolbar, click the + icon, with tooltip **Quick create** and select **Import YAML**.
 - d. In the Import YAML page, enter the manifests for a PVC, a deployment, and a service, which you already copied to your clipboard, and click menu:[Create].
 - e. You should see a page with the message **Resources successfully created**. Click on the **evalhubdb** deployment, the one with a "D" icon.
 - f. In its deployment details page, wait until you see a blue circle with a label of "1 Pod", indicating one pod from that deployment is running. Then scroll down to the **Conditions** table, and verify there is a row of type **Available** with status **True**.
3. Configure your EvalHub instance by creating a secret and an EvalHub resource.
- a. Open the [sample evalhub db secret](#) from the course samples repository with a web browser and open the manifest `scripts/evalhub-db-credentials.yaml`. Copy the entire contents of the manifest to the clipboard.

The full secret manifest is listed here, just for your reference:

```
apiVersion: v1
kind: Secret
metadata:
  name: evalhub-db-credentials
type: Opaque
stringData:
  db-url: "postgres://rhoai:dontdothisinprod@evalhubdb.redhat-ods-
applications.svc.cluster.local:5432/evalhub"
```

- b. In the left navigation pane, select **Workloads** > **Secrets** and verify that the selected project is still `redhat-ods-applications`. Then click **Create** > **From YAML**.
- c. In the **Create Secret** page, enter the manifests for a secret which you already copied to your clipboard, and click **Create**.
- d. In the **Secret details** page, you can scroll down and click **Reveal values** to see the database connection URL. It contains the database user name, password, and other connection parameters which match the database deployment you created in the previous step.
- e. In the left navigation pane, select **Home** > **API Explorer**. Then enter `EvalHb` in the search field. You will see only one resource, so click **EvalHub**.
- f. Open the [sample evalhub manifest](#) from the course samples repository with a web browser and open the manifest `scripts/evalhub.yaml`. Copy the entire contents of the manifest to the clipboard.



That manifest corresponds to the example from [2.3. Deploy EvalHub with the TrustyAI Operator](#) in the Red Hat OpenShift AI product documentation with a couple changes:

- Addition of **lighteval** to the list of providers
- The correct internal URL of the MLflow server

The full manifest is listed here, just for your reference:

```
apiVersion: trustyai.opendatahub.io/v1alpha1
kind: EvalHub
metadata:
  name: evalhub
spec:
  replicas: 1
  database:
    type: postgresql
    secret: evalhub-db-credentials
  providers:
    - lm-evaluation-harness
    - garak
    - guidellm
    - lighteval
  collections:
    - safety-and-fairness-v1
  env:
    - name: MLFLOW_TRACKING_URI
      value: "https://mlflow.redhat-ods-applications.svc.cluster.local:8443"
```

- In the EvalHub resource details page, select the **Instances** tab and verify that the selected project is still **redhat-ods-applications**. Then click **Create EvalHub**.
- In the YAML text editor, delete all of its sample contents and paste the evalhub manifest which you already copied to the clipboard.
- Click **Create** and, in the EvalHub details page, wait until the **Ready** condition is **True**.

Test your EvalHub server using the RHOAI dashboard

- Verify that EvalHub is available from the RHOAI dashboard.
 - Open your RHOAI dashboard. If it is already open, refresh your browser window so it picks changes to the dashboard pages and navigation.
 - In the left navigation pane, select **Develop & train > Evaluations**. The **Evaluations** page now displays the message **No existing evaluation runs**.
 - Click **Create Evaluation** and, on the **New evaluation run** page, click **Single benchmark**.
 - The **Single benchmark** page shows a long, paginated list of benchmarks offered by the providers configured on your EvalHub resource. Do *not* click any of them: Just seeing the list is sufficient to prove that your EvalHub server is functional.

Test your EvalHub server using a shell

- Verify that your EvalHub server is healthy.



The shell commands in this course assume a Bash shell on Fedora Linux. If your work machine runs Windows or macOS, you should either install a Bash shell or connect to the bastion host of your demo environment using SSH.

- a. Open a terminal and log in to the OpenShift cluster from your demo environment, as a cluster administrator, using the `oc login` command.
- b. Run the following commands to get the EvalHub URL and access its health endpoint. Your output should include the JSON field `status: healthy`.

```
$ oc project redhat-ods-applications
Now using project "redhat-ods-applications" on server "https://172.30.0.1:443".
$ EVALHUB_URL=https://$(oc get routes evalhub -o jsonpath='{.spec.host}')
$ curl -sk $EVALHUB_URL/api/v1/health
{"status":"healthy","timestamp":"2026-07-27T23:24:06.894598885Z","build":"0.3.0"}
```

5. Install the EvalHub CLI in a Python virtual environment.
 - a. Install the EvalHub CLI using `pip`. It should work with any version of Python from 3.12 on.

```
$ python --version
Python 3.14.6
$ python -m venv ~/venv/evalhub
$ source ~/venv/evalhub/bin/activate
... prompt changes ...
$ pip install "eval-hub-sdk[cli]"
...
Successfully installed annotated-types-0.8.0 anyio-4.14.2 certifi-2026.7.22
click-8.4.2 eval-hub-sdk-0.4.4 h11-0.16.0 httpcore-1.0.9
...
```

- b. Configure the EvalHub CLI to access your RHOAI instance.

Notice these commands refer to the `EVALHUB_URL` shell variable you set during the previous step.

```
$ evalhub config set base_url $EVALHUB_URL
Set 'base_url' in profile 'default'
$ evalhub config set tenant my-first-model
Set 'tenant' in profile 'default'
$ TOKEN=$(oc whoami -t)
$ evalhub config set token $TOKEN
Set 'token' in profile 'default'
```

- c. Verify that the EvalHub CLI can connect to the EvalHub server and list its known evaluation providers.

```
$ evalhub health
EvalHub service: healthy (640ms)
$ evalhub providers list
```

ID	NAME	DESCRIPTION
BENCHMARKS		
lm_evaluation_harness	LM Evaluation Harness	Comprehensive evaluation framework for language models with 180 benchmarks
lighteval	Lighteval	Lightweight LLM evaluation framework from Hugging Face
guidellm	GuideLLM	Performance benchmarking framework for LLM inference servers
garak	Garak	LLM vulnerability scanner and red-teaming framework

Summary

After completing these steps, you have verified that you have a minimal but functional EvalHub instance, which you can use to run evaluation jobs.

What's next

With both MLFlow and EvalHub deployed, the next chapter introduces model evaluation using EvalHub providers and benchmarks.

Model Evaluation with EvalHub

Goal

Evaluate LLMs using EvalHub providers, benchmarks, and collections on Red Hat OpenShift AI.



Not started yet.

Model Evaluation Concepts

Objective

Describe the evaluation providers and benchmarks available in EvalHub and explain how to manage evaluation jobs and interpret their results.



Not started yet.

Heading

Lorem ipsum

Another heading

Lorem ipsum

What's Next

Next, you run evaluation jobs on a self-hosted model and inspect the results.

Lab: Evaluate a Self-Hosted Model

Objective

Run evaluation jobs on a self-hosted model using EvalHub, monitor their progress, and inspect results and metrics from the RHOAI dashboard.



Pending review

Before you begin

You need access to the predeployed Red Hat OpenShift AI (RHOAI) instance provided in your lab environment. Ensure you have:

- Access credentials for the RHOAI dashboard
- A web browser to access the RHOAI dashboard
- Access credentials for the OpenShift web console

Also ensure that you already deployed MLflow and EvalHub by following the instructions from the

Instructions

In this lab, you will perform a quick evaluation run using the RHOAI dashboard and view its results in MLflow.

As you perform this activity, you will alternate between three windows:

- The Red Hat OpenShift AI (RHOAI) dashboard
- The Red Hat OpenShift web console
- A terminal window

Retrieve model name and endpoint, and label its project for TrustyAI

1. Find the name and endpoint of the model to evaluate.
 - a. Open the RHOAI dashboard and, in its left pane, select **AI hub** › **Models** and select the **Deployments** tab. There is a single row, named **llama-32-3b-instruct**.
 - b. Click its **Internal endpoint** link to raise a pop-up displaying its URL and click the **copy** icon to copy its URL to the clipboard.
 - c. Also take note of the model name, which is **llama-32-3b-instruct**, for use in the next step.
2. Label your project for TrustyAI network policies.
 - a. Open the OpenShift web console and select **Administration** › **Namespaces**. Enter **my** in the search field and click **my-first-model**.



You *cannot* change namespace labels by using **Home** › **Projects**. It seems to work, but the namespace resource associated with the project will remain unchanged.

- b. Click the **Actions** › **Edit labels** and, in the **Edit labels** popup, append **evalhub.trustyai.opendatahub.io/tenant=true**. Do not remove any of its other labels!
- c. Click **Save**.

Submit an evaluation job using the RHOAI dashboard

3. Find a quick benchmark to run.
 - a. Switch to the RHOAI dashboard and, in its left navigation pane, select **Develop & train** › **Evaluations**.
 - b. Click **Create new evaluation** and, on the **New evaluation run** page, click **Single benchmark**.
 - c. In **Single benchmark** page, select **Name** in the filter, and enter **quick** in the search field. You should see three benchmarks, each with a different label of **Safety**, **Performance**, and **Security**.
 - d. Click **Quick security smoke test** to open a right pane with details of that benchmark. Verify

that this benchmark comes from the **Garak** evaluation framework, which is unlikely to require large data sets from HuggingFaces.

4. Submit an evaluation job for the **Quick security smoke test** benchmark.
 - a. In the left pane, select **Develop & train > Evaluations** and, in the **Project** field, enter **my** and select **my-first-model**.
 - b. Click **Single benchmark**, select **Name** and enter **quick**. In the **Quick security smoke test** card, click **Use this benchmark**.
 - c. Enter **garak-quick** in the **Evaluation name** field and, in the **Create new experiment** field, enter **garak-evalhub**



There is an issue with the RHOAI dashboard which prevents it from retrieving names of existing experiments with MLflow. You must always create a new experiment, and ensure the name does not match any existing experiment in the same project nor the name of experiments that were marked for deletion. If you use the EvalHub CLI, you can either add to existing experiments or create new ones.

- d. Enter **llama-32-3b-instruct** in the **Model or agent name** field, and paste the URL you copied in the previous step in the **Endpoint URL** field. Append **/v1** to the URL so it is a valid OpenAI endpoint. The model is not configured to require authentication, so there is no need to provide an API key.
 - e. Click **Start evaluation run**. You should see a message **Evaluation "garak-quick" has been started**.



If there is any error, the error message might disappear before you can read it and it might seem that nothing happened. In that case, try starting the the evaluation run again, without making changes.

5. Monitor the evaluation job.
 - a. You should be still on the **Evaluations** page, which now displays a single row named **garak-quick** with a **Pending** status. Be patient while EvalHub downloads the container images required by the beבחmark.
 - b. Once the evaluation run starts, it should terminate very quickly, and you are likely to *not* see its **Running** status, but only see it transition to **Complete**. If it takes more than a few minutes without any change, refresh your RHOAI dashboard page.
 - c. Click **garak-quick** to review details of the evaluation run. That evaluation is expected to fail. Notice that its metrics indicate only one attack was attempted and it was successful. It's a common attack of "asking a model to do something it shouldn't do", which you would prevent, in production deployments, by enabling guard rails.
 - d. Optionally, you can switch to the OpenShift web console and enter **Workloads > Jobs** to monitor Kubernetes jobs in the **my-first-model** project. Notice that the name of the Job contains the ID of the EvalHub evaluation run, which unfortunately you *cannot* get from the RHOAI dashboard. You can also click **1 of 1** under **Completions** to view the Kubernetes pod managed by the job.

Submit an evaluation job using the CLI

6. Select a different benchmark to run.

- a. Open a shell, log in to OpenShift using the `oc` command, and activate the Python virtual environment you created in a [previous lab](#), then give to the EvalHub CLI your OpenShift authentication token.

```
$ source ~/venv/evalhub/bin/activate
output of the command
$ TOKEN=$(oc whoami -t)
$ evalhub config set token $TOKEN
Set 'token' in profile 'default'
```

- b. List the benchmarks from the `garak` provider.

```
$ evalhub providers describe garak
Provider: Garak
ID:      garak
Description: LLM vulnerability scanner and red-teaming framework
```

Benchmarks (8):

ID		NAME		CATEGORY	
METRICS					
owasp_llm_top10		OWASP LLM top 10 risk scan		security	
attack_success_rate					
avid		Full AI vulnerability scan		security	
attack_success_rate					
avid_security		AI security vulnerability scan		security	
attack_success_rate					
avid_ethics		AI ethics & bias scan		safety	
attack_success_rate					
avid_performance		AI performance degradation scan		performance	
attack_success_rate					
quality		Toxic & harmful content scan		safety	
attack_success_rate					
cwe		Software weakness exploitation scan		security	
attack_success_rate					
quick		Quick security smoke test		security	
attack_success_rate					

- c. Among these benchmarks, the **cwe** benchmark is expected to be quick to run.
7. Submit an evaluation job for the **cwe** benchmark.
 - a. Submit an evaluation job and save its ID in a shell variable.

```
$ evalhub eval run --name garak-cwe \
  --model-url http://llama-32-3b-instruct-predictor.my-first-
model.svc.cluster.local:8080/v1 \
  --model-name llama-32-3b-instruct \
  --provider garak --benchmark cwe \
  --experiment garak-evalhub
Job submitted: 11c1e53a-0086-4b22-9e5a-7c8ef0e07d50
$ ID=11c1e53a-0086-4b22-9e5a-7c8ef0e07d50
```

- b. Monitor the status of your second evaluation job. This time, it should transition to **running** very quickly because the Garak container images were already downloaded.

```
$ evalhub eval status
```

ID	NAME	STATE	PROVIDER
527d4579-0a75-482c-96a2-3c4b0222bb81	garak-quick	completed	garak
11c1e53a-0086-4b22-9e5a-7c8ef0e07d50	garak-cwe	running	garak

- c. if you prefer, refer to the evaluation ID to view its status unencumbered by other jobs. Repeat this command until it displays a state of **completed**, which shouldn't take more than a few seconds.

```
$ evalhub eval status $ID
Job:      11c1e53a-0086-4b22-9e5a-7c8ef0e07d50
Name:     garak-cwe
State:    completed
Model:    llama-32-3b-instruct (http://llama-32-3b-instruct-predictor.my-first-
model.svc.cluster.local:8080/v1)
Created:  2026-07-29 23:06:45.212450+00:00
```

Benchmarks (1):
cwe (garak): completed

Message: Evaluation job is completed

- d. Review the results of the evaluation run. Notice that the value of metric `attack_success_rate` is less than 0.3, which sets the `Pass` score you would see, for this evaluation run, in the RHOAI dashboard.

```
$ evalhub eval results $ID
```

BENCHMARK VALUE	PROVIDER	METRIC
cwe 0.1426	garak	attack_success_rate
cwe 0	garak	exploitation.JinjaTemplatePythonInjection_asr
cwe 0.1	garak	exploitation.SQLInjectionEcho_asr
cwe 0	garak	web_injection.ColabAIDataLeakage_asr
cwe 0	garak	web_injection.MarkdownImageExfil_asr
cwe 0.1641	garak	web_injection.MarkdownURIImageExfilExtended_asr
cwe 0.1758	garak	web_injection.MarkdownURINonImageExfilExtended_asr
cwe 0	garak	web_injection.MarkdownXSS_asr
cwe 0	garak	web_injection.PlaygroundMarkdownExfil_asr
cwe 0	garak	web_injection.StringAssemblyDataExfil_asr
cwe 0.0471	garak	web_injection.TaskXSS_asr

MLflow experiment: <https://mlflow.redhat-ods-applications.svc.cluster.local:8443/api/2.0/mlflow/experiments>

- e. If you need to inspect the Kubernetes job and pod associated with this evaluation run, to

perform troubleshooting, use the `job_id` label.

```
$ oc get job -l job_id=$ID -n my-first-model
NAME                                     STATUS
COMPLETIONS  DURATION  AGE
11c1e53a-0086-4b22-9e5a-7c-ed3b9527-d93a-42bd-9a7c-e2d913c21d17  Complete  1/1
2m20s      9m5s

$ oc get pod -l job_id=$ID -n my-first-model
NAME                                     READY  STATUS
RESTARTS  AGE
11c1e53a-0086-4b22-9e5a-7c-ed3b9527-d93a-42bd-9a7c-e2d913cx92gj  0/2
Completed  0      9m17s
```

View evaluation results on MLflow

8. View the two experiments from EvalHub in MLflow.
 - a. Return to the RHOAI dashboard and, in its left navigation pane, select **Develop & train > Experiments (MLflow)**. In the **Projects** field, enter `my` and select **my-first-model**.
 - b. The experiments page should display a single experiment. Click **garak-evalhub**.
 - c. In **Experiments** page for the **garak-evalhub** experiment, click **Evaluation runs**. It should display two runs, and their names match their EvalHub evaluation IDs. Click the one with a green ball, which is the most recent and corresponds to the second experiment, that got a **Pass** score.
 - d. That experiment contains no metrics, but select its **Overview** tab to visualize the same metrics you seen with the EvalHub CLI or with the **Evaluations** page of the RHOAI dashboard.

Summary

You performed two evaluation runs, using the Garak provider or evaluation framework: once using the RHOAI dashboard, and other using the EvalHub CLI. Then you visualized their results using both EvalHub and MLflow.

What's next

Having evaluated models with EvalHub, the next chapter focuses on evaluating agentic applications using the MLFlow SDK.

Agent Evaluation with MLFlow

Goal

Evaluate agentic applications using the MLFlow SDK with deterministic scorers and LLM judges on Red Hat OpenShift AI.



Not started yet.

Agentic Evaluation Concepts

Objective

Describe the different kinds of assessment for agentic applications and explain how to code MLFlow evaluations using scorers, judges, and data sets.



Not started yet.

Heading

Lorem ipsum

Another heading

Lorem ipsum

What's Next

Next, you evaluate a chat application using a deterministic scorer and view the results as MLFlow traces.

Lab: Evaluate an Agent with Deterministic Scoring

Objective

Starting from an agent which already produces MLflow traces, create a deterministic scorer, and run an MLFlow evaluation to verify response quality of that agent.



Pending review

Before you begin

You need access to the predeployed Red Hat OpenShift AI (RHOAI) instance provided in your lab environment. Ensure you have:

- Access credentials for the RHOAI dashboard
- A web browser to access the RHOAI dashboard
- Access credentials for the OpenShift web console

Also ensure that you already deployed MLflow and created a RHOAI workbench by following the instructions from [chapter 2](#).

Instructions

You will create and test an MLflow scorer functions, using a jupyter notebook. Then you will use the function for an MLflow evaluation of a simplistic agentic application that is provided to you.

As you perform this activity, you will alternate between two browser tabs:

- The Red Hat OpenShift AI (RHOAI) dashboard
- A Jupyter notebook from a RHOAI workbench

Review the sample agentic application

1. If you do not have a browser tab already open with the `test-mlflow` workbench, open it.
 - a. Open the RHOAI work bench and, in the left navigation pane, click **Projects**. Select **All projects** in the **Project** field. Then enter `my` in the **Filter by name** field and click **my-first-model** in the list of projects.
 - b. In the **Projects** page for `my-first-model`, select the **Workbenches** tab and click **test-mlflow** to open its Jupyter Lab in a new browser tab.
2. Import the provided agentic application from git and test it.
 - a. In the toolbar of left pane of the jupyter notebook, click the **Git Clone** icon and, in the **Clone a repo** pop-up, enter `https://github.com/RedHatQuickCourses/rhoai-mlflow-samples.git` which is the URL of the course samples repository. Click **Clone**. You should see, in the bottom right, the message `successfully cloned`.
 - b. In the left pane of the jupyter notebook, enter the `rhoai-mlflow-samples/evals` directory and double-click the `basic-agent.ipynb` notebook to open it in a notebook tab.
 - c. Review the Python code for the agent. It is a simplisic chat application which sets a hardcoded system prompt and submits a hardcoded question to a llama model, which is predeployed on your demo instance, using the OpenAI API.

The code in the notebook is split into multiple cells for readability and it could be concatenated into a single Python script to run in a shell, but then it would require changes to use external URLs to acess an inferencesservice and MLflow, and to use authentication tokens.



If you need more explanation about how that simplistic agent and its calls to MLflow tracing, please refer to the quick course about [MLflow traces](#).

- d. The notebook is ready to run as it is. In the top toolbar of the notebook tab, click the **fast-forward** icon with tooltip `Restart the kernel and run all cells`. In the **Restart Kernel** pop-up, click **Restart**.
- e. Be patient while the first cell downloads Python packages from the RHOAI index, and the second cell submits to the llama model a prompt asking what is RHEL. You can ignore any warnings about `IPProgress` and `ipywidgets`.

- f. If all goes well, you should get a reasonable and correct response, though a rather long one. Follows an example from testing these instructions:

****Red Hat Enterprise Linux (RHEL)****

- * Developed and maintained by Red Hat, Inc.
- * Licensed under the Red Hat Enterprise Linux License, which requires a subscription and support fees
- * Supports a wide range of hardware and software platforms
- * Includes commercial support, updates, and security patches for a fee
- * Typically used in production environments, data centers, and enterprise settings

****CentOS****

- * Originally based on RHEL 5, but now uses RHEL 8 as its base
- * Developed by a community of volunteers, led by Red Hat
- * Licensed under the GNU General Public License (GPL), which is free and open-source
- * No subscription or support fees required
- * Often used in development, testing, and education environments, as well as in some production environments where cost is a concern

In summary, RHEL is a commercial Linux distribution that requires a subscription and support fees, while CentOS is a free and open-source distribution that is based on RHEL but is not directly supported by Red Hat.

3. Visualize the MLflow trace generated by the agent.

- Switch to the RHOAI dashboard and, in the left navigation pane, select **Develop & train > Experiments (MLflow)**. If the current project is not already **my-first-model** enter **my** in the **Project** field and click **my-first-model**.
- The **Experiments** page should display both the **garak-evalhub** experiment from the previous chapter and the **simple-agent** experiment from this activity. Click **simple-agent** to open its experiment details page and click **Traces**.
- You should see a single row, which displays the request and response from the agent, and metrics such as token count and execution time, but no assessments. The **OK** state just means the agent performed the request and produced a response without execution errors.
- Click either the trace id or the request on the single row to open a trace summary pane. Then click **Show assessments** to open its assessments pane. It is empty, displaying only buttons to add feedback and expectations.
- Return to the traces table by clicking the close button of the trace summary pane, which is the **x** icon at its top right corner. Do *not* click the similar icon inside the assessments pane!

Evaluate the agent using a deterministic scorer function

4. Create and test a deterministic scorer function to assess the length of a response.



You can see the complete code, ready to run evaluations, in the [deterministic-eval-complete.ipynb](#) notebook, in the [course samples repository](#).

- a. Copy the basic-agent notebook to a new file which you will edit.

Switch to your jupyter labs and, in its left pane, right click **basic-agent.ipynb** and select **Duplicate**. Then right-click the new file **basic-agent-Copy1.ipynb** and select **Rename**. Enter **deterministic-eval.ipynb** as the new file name. Finally, double-click **deterministic-eval.ipynb** to open the notebook in a new editor tab.

- b. Enter, inside the cell, the code for a scorer function which checks the length of a response based on the number of words.

In the **deterministic-eval.ipynb** notebook, select the last cell, having the comment **###Cell 6: run the agent**. Then click, in the cell's toolbar, the icon to **Insert a cell below**. Inside the new cell, enter the following code.

```
###Cell 7: deterministic scorer

from mlflow.genai.scorers import scorer
from mlflow.entities import Feedback

@scorer
def one_to_three_short_sentences(outputs: str) -> Feedback:
    """15-20 words per sentence, up to three sentences"""
    word_count = len(outputs.split())
    metric={
        "word-count": word_count,
        "min": 20,
        "max": 60
    }
    pass_fail = word_count >= 20 and word_count <= 60
    if word_count < metric["min"]:
        return Feedback(
            name="short_response",
            value=False,
            rationale=f"Total word count was lower than {metric['min']}: {word_count}.",
            metadata=metric
        )
    elif word_count > metric["max"]:
        return Feedback(
            name="short_response",
            value=False,
            rationale=f"Total word count was higher than {metric['max']}: {word_count}.",
```

```

        metadata=metric
    )
else:
    return Feedback(
        name="short_response",
        value=True,
        rationale=f"Total world count was between {metric["min"]} and
{metric["max"]}: {word_count}.",
        metadata=metric
    )

```

- c. Click, in the cell's toolbar, the **play** icon with tooltip **Run this cell and advance** to define the function and create a new notebook cell. Running the cell produces no output.
- d. Inside the new cell, enter the code to test the scorer function from hardcoded examples.

```

###Cell 8: test scorer

feedback = one_to_three_short_sentences(
    outputs="RHEL (Red Hat Enterprise Linux) is a Linux distribution."
)
print(f"Test scorer with a too short response: {feedback.value} -
{feedback.rationale}")

feedback = one_to_three_short_sentences(
    outputs="RHEL (Red Hat Enterprise Linux) is a commercial Linux distribution,
supported by Red Hat and produced from CentOS Stream and Fedora Linux."
)
print(f"Test scorer with an OK long response: {feedback.value} -
{feedback.rationale}")

feedback = one_to_three_short_sentences(
    outputs="""
        Red Hat Enterprise Linux (RHEL) is a commercial, stable, and secure
version of Linux, supported by Red Hat and produced from from multiple
upstreams.
        The direct upstream to RHEL is the CentOS Stream Linux distribution,
which is a free, community-supported distribution.
        The Fedora Linux is a direct upstream to CentOS, and thus an indirect
upstream to RHEL.
        RHEL is suitable for production environments, whereas Fedora is more
experimental and intended for developers and power users.
        Fedora is often used as a testing ground for new features and
technologies before they are included in RHEL.
        Fedora Linux is free and community-supported.
    """
)
print(f"Test scorer with a too long response: {feedback.value} -
{feedback.rationale}")

```

- e. Click, in the cell's toolbar, the **play** icon with tooltip **Run this cell and advance** to run the test code. Running the cell produces the following output.

```
Test scorer with a too short response: False - Total world count was lower than 20: 9.
Test scorer with an OK long response: True - Total world count was between 20 and 60: 22.
Test scorer with a too long response: False - Total world count was higher than 60: 100.
```

5. Create code for running an MLflow evaluation using the deterministic scorer.

- a. In the new, empty cell, created by running the test code for the scorer function, enter the code to define an evaluation data set.

```
###Cell 9: hardcoded evaluation data set

eval_dataset = [
    {
        "inputs": {"question":
            "What's the difference between RHEL and CentOS?"},
        "expectations": {"expected_response":
            """
            RHEL and CentOS are both Linux distributions. RHEL is a proprietary
            distribution from Red Hat, while CentOS is community-supported. CentOS is a
            direct upstream of RHEL.
            """},
    },
    {
        "inputs": {"question":
            "What's the difference between RHEL and Fedora Linux?"},
        "expectations": {"expected_response":
            """
            RHEL and Fedora are both Linux distributions. RHEL is a proprietary
            distribution from Red Hat, while Fedora is community-supported. Fedora is an
            upstream of RHEL.
            """},
    },
]
```

- b. Click, in the cell's toolbar, the **play** icon with tooltip **Run this cell and advance** to define the array and create a new notebook cell. Running the cell produces no output.
- c. Inside the new cell, enter the code to wrap the agent inside a predictor function and run an MLflow evaluation using your scorer function.



The `MLFLOW_GENAI_EVAL_SKIP_TRACE_VALIDATION` environment variable prevents MLflow from performing a canary test using the first row of the evaluation data, which would produce duplicate traces and assessments for

it.

```
###Cell 10: run evaluation

def my_predict_fn(question: str) -> str:
    return agent("no token", system_instructions, input_text)

os.environ["MLFLOW_GENAI_EVAL_SKIP_TRACE_VALIDATION"]="true"

try:
    results = mlflow.genai.evaluate(
        data=eval_dataset,
        predict_fn=my_predict_fn,
        scorers=[one_to_three_short_sentences],
    )
except Exception as e:
    print(f"An error occurred: {e}")
```

- d. Click, in the cell's toolbar, the **play** icon with tooltip **Run this cell and advance** to run the evaluation. Monitor the evaluation progress meter, and wait until it reaches 100%, meaning it run all scorers over the complete evaluation data set. When done, you should see an output similar to the following.

```
Evaluating: 100%|██████████| 2/2 [Elapsed: 00:01, Remaining: 00:00] [predict_fn:
98%, scorers: 2%]
```

6. View evaluation results on MLflow traces

- a. Switch to your RHOAI dashboard, which should already be in the details page for the **simple-agent** experiment. If it is not, select **Develop & train > Experiments (MLflow)**, in the left navigation pane, and, in the **Project** field, enter **my** and select **my-first-model** and click **simple-agent**.
- b. In the **simple-agent** experiment details page, click **Traces** and notice, besides the two new traces, the new columns **short_response** which display results from the scorer function. If you do not see these changes, click the **Refresh** icon next to the **Time range** field.

Also notice that the trace for the previous run of the agent, which did no evaluation, displays **null** for that column.

Finally, notice that the top of the table displays a summary of the **short_response** score for all traces, which a 100% for the **False** value.

- c. Click the trace ID or the request of the most recent trace, which is the first row of the traces table, to view its trace summary panel. It should display an **Assessment** pane, with the value of the **short_response** scorer. Click the **>** icon to expand the feedback results and view the rationale for the result.
- d. Return to the traces table by clicking the close button of the trace summary pane pane,

which is the **x** icon at its top right corner. Do *not* click the similar icon in the assessments pane!

Improve the agent with a new system prompt from the MLflow prompt registry

7. Record a new system prompt in MLflow.

You already seem that the llama model produces very long responses, by default. It is standard practice to tweak the system prompt to control the tone and format of outputs from the agent. But, instead of hardcoding a new system prompt in the agent, retrieve the prompt from MLflow.

- Still in the **simple-agent** experiment details page, click **Prompts**. Then click **Create prompt**.
- In the **Create prompt** pop-up, enter **short-responses** in the **Name** field, and enter the following in the **Prompt** field:

You are a helpful assistant which provides short and objective responses, composed of one to three short sentences.

- Click **Create** to save the prompt. Notice that it is saved as **version 1**.

8. Change the agent to use the new system prompt.



You can see the complete code, ready to use the new system prompt, in the **deterministic-eval-new-prompt.ipynb** notebook, in the [course samples repository](#).

- Switch to your jupyter lab and its editor tab with the **deterministic-eval.ipynb** notebook and scroll up to its sixth cell, that starts with **###Cell 6: run the agent**.
- Make the following changes to the code inside the cell to load and format the prompt from MLflow and use it as the system prompt. Even though our new prompt contains no placeholders, it must still be formatted before passing to an LLM.

```
...
try:
    system_prompt = mlflow.genai.load_prompt("prompts:/short-responses/1")
    final_prompt = system_prompt.format()
    response = agent("no token", final_prompt, input_text)
    print(response)
except Exception as e:
    print(f"An error occurred: {e}")
```

- In the cell's toolbar, click the **play** icon with tooltip **Run this cell and advance** to run the agent using the new system prompt. Its output should be very different from its previous run, and similar to the following.

RHEL (Red Hat Enterprise Linux) and CentOS are both Linux distributions, but they differ in their licensing and support. RHEL is a commercial distribution,

while CentOS is an open-source distribution based on RHEL.

9. Run evaluations using the new system prompt

- Scroll down to the tenth cell of the notebook, that starts with `###Cell 10: run evaluation`, and change its code to use the new system prompt.

```
...
def my_predict_fn(question: str) -> str:
    system_prompt = mlflow.genai.load_prompt("prompts:/short-responses/1")
    final_prompt = system_prompt.format()
    return agent("no token",final_prompt, input_text)
...
```

- In the cell's toolbar, click the **play** icon with tooltip `Run this cell and advance` to run the agent using the new system prompt. The output is the same as the previous attempt, because the code for evaluations does not print anything about the results, but just a progress meter.
- Switch to the RHOAI dashboard, which should still be in the **Prompts** tab of the experiment details page for `simple-agent`. Select **Traces** to open its list of traces.
- You should see the new traces, from running the agent and the evaluation a second time. If you do not, click the **Refresh** icon near the **Time range** field.

The new traces should display the values `null` and `True` in the `short_response` column, proving that the new system prompt had the desired effect. The summary at the top of the table should now split 50%/50% between `True` and `False`.

Summary

You created a scorer functions to assess the length of responses produced by the model, then you run an evaluation which showed the responses were too long. Then you crafted a new system prompt, which you stored in the MLflow prompt registry, to direct the model to produce shorter responses. After testing the agent using the new system prompt, you run the same evaluation again, proving that now responses were are short, as desired.

What's next

Next, you evaluate the same application using built-in and custom LLM judges for more nuanced assessment.

Lab: Evaluate an Agent with an LLM Judge

Objective

Perform evaluations using built-in and custom LLM judges with self-hosted models, and inspect evaluation results through MLFlow traces.



Work In Progress

Before you begin

You need access to the predeployed Red Hat OpenShift AI (RHOAI) instance provided in your lab environment. Ensure you have:

- Access credentials for the RHOAI dashboard
- A web browser to access the RHOAI dashboard
- Access credentials for the OpenShift web console

Also ensure that you already deployed MLflow and created a RHOAI workbench by following the instructions from [chapter 2](#), and that you already cloned the sample notebooks from the course samples repository, by following the instructions from the [previous lab](#).

Instructions

You will create and test two LLM judges, one provided by MLflow and another using your custom prompt. You will use them together for an MLflow evaluation of a simplistic agentic application that is provided to you

As you perform this activity, you will alternate between two browser tabs:

- The Red Hat OpenShift AI (RHOAI) dashboard
- The Jupyter Lab from a RHOAI workbench

Create and test a LLM judge using one of the judges provided by MLflow

1. If you do not have a browser tab already open with the `test-mlflow` workbench, open it and, in its Jupyter Lab, enter the `rhoai-mlflow-samples/evals` directory.
 - a. Open the RHOAI work bench and, in the left navigation pane, click **Projects**. Select **All projects** in the **Project** field. Then enter `my` in the **Filter by name** field and click **my-first-model** in the list of projects.
 - b. In the **Projects** page for `my-first-model`, select the **Workbenches** tab and click `test-mlflow` to open its Jupyter Lab in a new browser tab.
 - c. Copy the `simple-agent` notebook to a new file which you will edit. In the left pane of the jupyter notebook, enter the `rhoai-mlflow-samples/evals` directory, right click `simple-agent.ipynb`, and select **Duplicate**. Then right-click the new file `simple-agent-Copy1.ipynb` and select **Rename**. Enter `llm-eval.ipynb` as the new file name. Finally, double-click `llm-eval.ipynb` to open the notebook in a new editor tab.
 - d. In the toolbar of the notebook, click the **fast-forward icon** with tooltip `Restart the kernel and run all cells`. In the **Restart kernel** pop-up, click **Restart**. The output of the last cell should provide a short explanation about the differences between RHEL and CentOS.



the code uses the prompt added to the MLflow prompt registry during the [previous lab](#).

2. Create and test an LLM judge to assess the correctness of a response.

You can see the complete code, ready to use the built-in LLM judge, in the [llm-eval-builtin.ipynb](#) notebook, in the course samples repository.

- a. Enter the model configurations for MLflow LLM judges.

In the `llm-eval.ipynb` notebook, select the third cell, with the comment `###Cell 3: environment settings` and append configuration and environment variables to configure MLflow to use the llama model available in your demo environment to run LLM judges.

```
###Cell 3: environment settings

import os

BASE_URL = "http://llama-32-3b-instruct-predictor.my-first-
model.svc.cluster.local:8080/v1"
model_name = "llama-32-3b-instruct"

MLFLOW_TRACKING_URL = "https://mlflow.redhat-ods-
applications.svc.cluster.local:8443"
MLFLOW_EXPERIMENT = "simple-agent"

os.environ["MLFLOW_TRACKING_AUTH"]="kubernetes-namespaced"

# New code starts here
JUDGE_URL = BASE_URL
judge_model = model_name
os.environ["OPENAI_API_BASE"] = JUDGE_URL
os.environ["OPENAI_API_KEY"] = "no token"
```

- b. Click, in the cell's toolbar, the **play** icon with tooltip `Run this cell and advance` to define the new variables. Running the cell produces no output.
- c. Enter the code to instantiate a `Correctness` judge from MLflow.

In the `llm-eval.ipynb` notebook, select the last cell, having the comment `###Cell 6: run the agent`. Then click, in the cell's toolbar, the icon to **Insert a cell below**. Inside the new cell, enter the following code.

```
###Cell 7: LLM judges

from mlflow.genai.scorers import Correctness
from mlflow.genai.judges import make_judge

correctness_judge = Correctness(
    model=f"openai://{judge_model}"
)
```

- d. Click, in the cell's toolbar, the **play** icon with tooltip `Run this cell and advance` to define the function and create a new notebook cell. Running the cell produces no output.

- e. Inside the new cell, enter the code to test the LLM judge from hardcoded examples.

```
###Cell 8: test judges

print ("Test built-in judge:\n")
feedback = correctness_judge(
    inputs={"question": "What is RHEL?"},
    outputs={"response": "RHEL is a Linux distribution from Red Hat."},
    expectations={"expected_response": "RHEL (Red Hat Enterprise Linux) is a Linux distribution."},
)
print(f"▯ Test judge with correct definition of RHEL: {feedback.value}\n{feedback.rationale}")

feedback = correctness_judge(
    inputs={"question": "What is RHEL?"},
    outputs={"response": "RHEL (Rebooting Helps Every Lag) is pun on IT struggles."},
    expectations={"expected_response": "RHEL (Red Hat Enterprise Linux) is a Linux distribution."},
)
```

- f. Click, in the cell's toolbar, the **play** icon with tooltip **Run this cell and advance** to run the test code. Running the cell produces an output similar to the following.

Test built-in judge:

▯ Test judge with correct definition of RHEL: yes

The first statement in the claim is 'RHEL' which is supported by the document as the question and response both mention 'RHEL'. The second statement in the claim is 'is a Linux distribution' which is also supported by the document as the response explicitly states 'a Linux distribution from Red Hat'.

▯ Test judge with incorrect definition of RHEL: no

The claim states RHEL is a Linux distribution. The document does not mention Linux distribution, it mentions a pun on IT struggles.

Create and test a LLM judge using a custom prompt

3. Create and test an LLM judge to assess that a response contains no acronyms

You can see the complete code, ready to use the built-in LLM judge, in the [llm-eval-custom.ipynb](#) notebook, in the course samples repository.

- a. Select the first cell that you just added to the notebook, the one starting with **###Cell 7: LLM judges** and append code to create a custom judge.

```
###Cell 7: LLM judges
```

```

from mlflow.genai.scorers import Correctness

correctness_judge = Correctness(
    model=f"openai:{judge_model}"
)

# New code starts here
from mlflow.genai.judges import make_judge

acronyms_judge = make_judge(
    name="expanded_acronyms",
    instructions=(
        "Check if this text expands any common IT acronyms or product names:\n"
        "{{ outputs }}"
        "\n\nEVALUATION STEPS:\n"
        "1. Scan the text for IT acronyms (e.g., TCP, IP, AI, LLM) or similar names (e.g., Vim, GNOME, LAME, RHEL).\n"
        "2. Check if the text explicitly expands them to their full words (e.g., 'LLM (Large-Language Model)', 'GNOME (GNU Network Object Model Environment).\n"
        "3. Merely explaining or describing a concept (e.g., 'Vim is a text editor') (e.g. 'RHEL is an operating system') does NOT count as an expansion.\n\n"
        "BOOLEAN OUTPUT:\n"
        "- Return True if you found ONE OR MORE acronym expansions.\n"
        "- Return False if the response is completely FREE of acronym expansions (NO expansions found).\n"
    ),
    feedback_value_type=bool,
    model=f"openai:{judge_model}",
    inference_params={"temperature": 0.0},
    base_url=f"{JUDGE_URL}/chat/completions"
)

```

- b. Click, in the cell's toolbar, the **play** icon with tooltip **Run this cell and advance** to define the function and create a new notebook cell. Running the cell produces no output.
- c. Select the second cell that you just added to the notebook, the one starting with **###Cell 8: test judges** and append code to test the custom judge.

```

###Cell 8: test judges

print ("Test built-in judge:\n")
...

# New code starts here
print ("\nTest LLM judge:\n")

feedback = acronyms_judge(
    inputs={"question": "What is RHEL?"},
    outputs={"response": "RHEL is a Linux distribution."},

```

```

)
print(f"▯ Test judge with RHEL (no expansion): {feedback.value}\n
{feedback.rationale}")

feedback = acronyms_judge(
    inputs={"question": "Who supports RHEL?"},
    outputs={"response": "Red Hat provides commercial support for users of
Linux."},
)
print(f"▯ Test judge with RHEL (no expansion): {feedback.value}\n
{feedback.rationale}")

feedback = acronyms_judge(
    inputs={"question": "What is RHEL?"},
    outputs={"response": "RHEL (Red Hat Enterprise Linux) is a Linux
distribution."},
)
print(f"▯ Test judge with RHEL (expanded)': {feedback.value}\n
{feedback.rationale}")

```

- d. Click, in the cell's toolbar, the **play** icon with tooltip **Run this cell and advance** to run the test code. Running the cell produces an output similar to the following.

Test built-in judge:

...

Test LLM judge:

▯ Test judge with RHEL (no expansion): False

The text does not explicitly expand any IT acronyms or product names. Although it mentions 'RHEL', it only provides a brief description without expanding the acronym to its full word 'Linux distribution'. The text does not meet the criteria for an acronym expansion.

▯ Test judge with RHEL (no expansion): False

The text does not explicitly expand any IT acronyms or product names. It only mentions 'Linux', which is a common operating system, but does not provide any expansion or explanation of the acronym.

▯ Test judge with RHEL (expanded)': True

The text explicitly expands the acronym 'RHEL' to its full word 'Red Hat Enterprise Linux', indicating that it meets the criteria for evaluating the performance of an AI agent on this query.

Run an MLflow evaluation using both LLM judges

4. Create code for running an MLflow evaluation using the two LLM judges you just tested.
 - a. In the new, empty cell, created by running the test code for the custom judge function, enter the code to define an evaluation data set.

```

###Cell 9: hardcoded evaluation data set

eval_dataset = [
    {
        "inputs": {"question":
            "What's the difference between RHEL and CentOS?"},
        "expectations": {"expected_response":
            """
            RHEL and CentOS are both Linux distributions.
            RHEL is a commercial distribution from Red Hat, while CentOS is a
            free distribution.
            CentOS Stream is a direct upstream of RHEL.
            """},
    },
    {
        "inputs": {"question":
            "What's the difference between RHEL and Fedora Linux?"},
        "expectations": {"expected_response":
            """
            RHEL and Fedora are both Linux distributions.
            RHEL is a commercial distribution from Red Hat, while Fedora is a
            free distribution.
            Fedora is an upstream of RHEL.
            """},
    },
]

```

- b. Click, in the cell's toolbar, the **play** icon with tooltip **Run this cell and advance** to define the array and create a new notebook cell. Running the cell produces no output.
- c. Inside the new cell, enter the code to wrap the agent inside a predictor function and run an MLflow evaluation using your scorer function. If you need, follow the instructions there to store the prompt.

```

###Cell 10: run evaluation

def my_predict_fn(question: str) -> str:
    system_prompt = mlflow.genai.load_prompt("prompts:/short-responses/1")
    final_prompt = system_prompt.format()
    return agent("no token", final_prompt, input_text)

os.environ["MLFLOW_GENAI_EVAL_SKIP_TRACE_VALIDATION"]="true"

try:
    results = mlflow.genai.evaluate(
        data=eval_dataset,
        predict_fn=my_predict_fn,
        scorers=[correctness_judge, acronyms_judge],
    )

```



```
except Exception as e:
    print(f"An error occurred: {e}")
```

- d. Click, in the cell's toolbar, the **play** icon with tooltip **Run this cell and advance** to run the evaluation. Monitor the evaluation progress meter, and wait until it reaches 100%, meaning it run all scorers over the complete evaluation data set. When done, you should see an output similar to the following.

```
2026/08/03 16:47:51 INFO mlflow.models.evaluation.utils.trace: Auto tracing is
temporarily enabled during the model evaluation for computing some metrics and
debugging. To disable tracing, call 'mlflow.autolog(disable=True)'.
Evaluating: 100%|██████████| 2/2 [Elapsed: 00:08, Remaining: 00:00] [predict_fn:
22%, scorers: 78%]
```

5. View evaluation results on MLflow traces

- a. Switch to your RHOAI dashboard, which should already be in the details page for the **simple-agent** experiment. If it is not, select **Develop & train > Experiments (MLflow)**, in the left navigation pane, and, in the **Project** field, enter **my** and select **my-first-model** and click **simple-agent**.
- b. In the **simple-agent** experiment details page, click **Traces** and notice, besides a few new traces, the new columns **Correctdness** and **expanded_acronyms**. If you do not see these changes, click the **Refresh** icon next to the **Time range** field.
- c. Consider just the two most recent traces, which contains the results of the last MLflow evaluation run on a sample of two prompts, and mouse over the values on their **Correctdness** column to view the reasoning for the **Fail** and **Pass** grades of each trace.
- d. On the same traces, mouse over the values of their **expanded_acronyms** column to view the reasoning between their **True** or **False** scores.
- e. With any of the traces, if you cannot read the full reasoning by using mouse over, click on the trace id or its request and, if the **Assessments** page is not shown, click **Show assessments**. Then you can expand each assessment to fully visualize its reasoning, and you can scroll down and up as you need.
- f. You may note that the **expanded_acronyms** custom judge is not very reliable, which might be because of the quality of the prompt or by limitations of the llama model. It is frequent to use different models or even larger models for judges than the models that might be good enough for an application.
- g. You may get assessments and reasonings which seem odd or incorrect. It is the nature of LLMs that sometimes they hallucinate. Feel free to run the evaluation multiple times, compare the responses of each attempt, and how the LLM judges assess each one.

Summary

You two LLM judges, one to verify if an response is correct, by comparing with a correct response provided in the evaluation data set, and another to verify if a response expands acronyms, which would make it longer than necessary. After testing each judge, you combined both in a single

evaluation again, demonstrating that you can use multiple judges on the same evaluation data and visualize the results of multiple judges in the same MLflow trace.

What's next

Congratulations on completing this course! You are now able to deploy MLflow and EvalHub on Red Hat OpenShift AI, evaluate both models and agentic applications, and visualize results using the RHOAI dashboard.